

Communication

- connection
 - between processors
 - between processors and memory
- alternatives
 - symmetric shared memory systems - bus
 - distributed shared memory systems – interconnection networks

Bus Communication

- economic
- simple design
- lower performance
- does not scale well with the number of processors
- very important - cache memories
 - coherence problems

Interconnection Networks (1)

- between
 - processors
 - processors and memory
- goals
 - flexibility
 - performance
 - more accesses in parallel

Interconnection Networks (2)

- types of connection
 - total - each one to all others
 - partial - some pairs of components are not directly connected
- processor-memory connection
 - more memory circuits
 - can be accessed in parallel by different processors

Memory Coherence

- all processors must use the last value written into a shared variable
- the problem - cache memories
- goal - every shared variable has the same value
 - in all caches
 - in the main memory

Coherence - Write-back Cache

x - shared variable

Processor	Action	Cache A value	Cache B value	Memory value
				9
A	i=x;	9		9
B	j=x;	9	9	9
A	x=5;	5	9	9
B	k=x;	5	9	9

Coherence - Write-through Cache

x - shared variable

Processor	Action	Cache A value	Cache B value	Memory value
				9
A	i=x;	9		9
B	j=x;	9	9	9
A	x=5;	5	9	5
B	k=x;	5	9	5

What Is Coherence? (1)

1. execution order

- a) processor P writes to variable X
- b) then processor P reads X
 - there are no writes to X between a) and b)

→ read b) returns the value written by a)

What Is Coherence? (2)

2. coherent vision of the memory
 - a) processor P writes to variable X
 - b) then processor Q ($Q \neq P$) reads X
 - there are no writes to X between a) and b)
 - enough time passed between a) and b)
- read b) returns the value written by a)

What Is Coherence? (3)

3. write serialization

- a) processor P writes to variable X
- b) processor Q ($Q=P$ or $Q \neq P$) writes to variable X

→ all processors see the two writes in the same order

- not necessarily a) before b)

Maintaining Cache Coherence

- coherence maintenance protocols
- based on the information about the cache lines
 - *invalid* - data is not valid
 - *dirty* - only the current cache has the correct (updated) value
 - *shared* - the current cache has the correct (updated) value, so do main memory and possibly other caches

Types of Protocols

- *directory based*
 - information about each cache line - kept in a single place
- *snooping*
 - each cache has a copy of the shared line
 - no centralized information
 - caches "listen" to the bus
 - detect changes brought to cache lines

Cache Update

- each cache announces the changes it makes
- the other caches react
- only write operations matter
- variants
 - write invalidate
 - write update

Write Invalidate (1)

- a processor changes the value of a location
- the change is made in its own cache
 - all other caches are notified
- every other cache
 - has no copy of that location - no action
 - has a copy of that location - invalidates its corresponding line
 - the correct value will be requested when needed

Write Invalidate (2)

x - shared variable

Proc.	Action	Cache reaction	Cache A	Cache B	Memory
					9
A	i=x;	read miss	9		9
B	j=x;	read miss	9	9	9
A	x=5;	invalidation	5	inv.	9
B	k=x;	read miss	5	5	5

Write Update (1)

- a processor changes the value of a location
- the change is made in its own cache
 - all other caches are notified
 - the new value is broadcasted
- every other cache
 - has no copy of that location - no action
 - has a copy of that location - gets the new value

Write Update (2)

x - shared variable

Proc.	Action	Cache reaction	Cache A	Cache B	Memory
					9
A	i=x;	read miss	9		9
B	j=x;	read miss	9	9	9
A	x=5;	invalidation	5	5	5
B	k=x;	read hit	5	5	5

Write Invalidate vs. Write Update

(1)

- multiple successive writes to the same location
 - write invalidation - only one invalidation (first time)
 - write update - an update for each write
 - more favorable - write invalidation

Write Invalidate vs. Write Update

(2)

- multiple writes to the same cache line
 - changing a location requires invalidating/updating the whole line
 - write invalidation - only one invalidation (first time)
 - write update - write update - an update for each write
 - more favorable - write invalidation

Write Invalidate vs. Write Update

(3)

- "response time"
 - the time between writing a value by a processor and reading that value by another processor
 - write invalidation - first invalidation, then read (when necessary)
 - write update - immediate update
 - more favorable - write update

Write Invalidate vs. Write Update

(4)

- both variants have advantages and drawbacks
- write invalidate - (much) lower usage of memory and buses
- write update - higher cache hit rate
- more often used - write invalidate

III. Peripheral Devices

Peripheral Device

- provides a certain kind of communication
 - between the processor and the "outside world"
- to manage communication with the processor, it includes an I/O controller
 - which usually contains a series of registers that store information necessary for communication
 - data
 - state information
 - commands (from the processor)

Input-output (I/O)

- how the system sees communication
 - memory-mapped I/O
 - read/write operations are seen as though they are performed on memory locations
 - I/O addresses - within the memory address space
 - the same control signals as for memory
 - isolated I/O
 - I/O addresses - separate from memory addresses
 - control signals - different from those of memory

Communication Modes (1)

- programmed I/O
 - the program waits in a loop until the peripheral device initiates a transfer
 - efficient if the moment when the device will request communication is known in advance
 - useless consumption of processor time

Communication Modes (2)

- Direct Memory Access (DMA)
 - a specialized controller (DMA controller) manages the transfer
 - very fast
 - it takes control of the buses and transfers data directly between the device and memory
 - the processor is bypassed
 - useful for transferring large amounts of data

Communication Modes (3)

- interrupt-driven I/O
 - when a peripheral device wishes to communicate, it notifies the processor
 - through an interrupt request
 - during the rest of the time, the processor may carry out other tasks
 - the most flexible method

The Buses (1)

- communication paths for the information
- a bus is a unique connection between more than 2 components
- bus description
 - electrical signals
 - communication rules - must be respected by all parts involved
 - connecting mode

The Buses (2)

Bus access

- more entities may request access simultaneously
- an arbitration procedure is required
 - decides who will be granted access
 - the other entities must wait until the bus becomes free again

Bus Arbitration

Types of bus arbitration

- centralized
 - the decision is made by a dedicated circuit (arbiter)
- de-centralized
 - components negotiate to each other
 - based on the rules that define the way the bus works

Connecting to the Bus

- electrical issues
- more circuits connected together
 - input and output
- cannot connect more outputs
 - different voltage levels would destroy the circuits
- one solution - multiplexing
 - all outputs are connected to a multiplexer

Tri-state Circuits

- output has 3 possible states
 - 0
 - 1
 - high impedance (*High-Z*)
- first two states correspond to usual values
- third state means decoupling from the bus
 - just as the output of the circuit would not be connected to the bus

Open-collector Circuits

- in some cases are called *open-drain*
 - depends on technology used
- it is possible to connect more outputs together
- result value - Boolean AND between the outputs that are connected

Buses - General View

Advantages

- bus activity - easy to control
- economic - relatively simple structure

Drawback

- lower performance
 - only 2 components can communicate at a certain moment

IV. The Interrupt System

What Is an Interrupt?

- the processor can suspend the execution of the current program
- goal - handling unexpected situations
- after that, the interrupted program is resumed
- initial purpose - communicate to peripherals
- the processor does not "wait" for peripherals
 - they notify the processor when necessary

Hardware Interrupts


- maskable
 - the processor may refuse to service them
 - depends on the value of the IF (*Interrupt Flag*)
flip-flop: 1 - accept, 0 - deny
 - IF can be modified by software
- non-maskable
 - the processor always services them

The Interrupt Controller (1)

- specialized circuit
- collects interrupt requests from the peripherals
- sends them to the processor
- arbitrates the conflicts (more requests coming simultaneously)
 - each peripheral is assigned a certain priority

The Interrupt Controller (2)

- initially - Intel 8259
 - multiple controllers can be used (cascading)



- today - integrated into the chipset

Interrupt Handling - Phases (1)

- the peripheral generates an interrupt request on its IRQ_i line
- the controller activates signal INT
- the processor checks the value of the IF flip-flop
 - only for maskable interrupts
 - if 0 - refuses the request; stop
 - if 1 - responds by activating signal INTA

Interrupt Handling - Phases (2)

- execution of the current program is suspended
- the registers (including the PC) are saved on the stack
- the IF flip-flop is reset
 - blocks the execution of another interrupt while the current interrupt is serviced
 - can be set again by software

Interrupt Handling - Phases (3)

- identify the peripheral (the source of the request)
 - the controller places a *type* byte on the data bus
 - it indicates the peripheral that made the request
 - at most $2^8=256$ interrupt sources
 - each source has its own service routine
 - different peripherals - different services

Interrupt Handling - Phases (4)

- determine the address of the service routine
 - physical address 0 - interrupt vector table
 - contains the addresses of all service routines
 - size: $256 \text{ addresses} \times 4 \text{ bytes} = 1 \text{ KB}$
 - type byte = $n \rightarrow$ the address of the service routine = $n \times 4$

Interrupt Handling - Phases (5)

- jump to the address of the service routine
- execute the service routine
- go back to the interrupted program
 - restore the value of the IF flip-flop
 - restore the values of the registers (from stack)
 - resume the execution of the program from the point it was interrupted

Extension

- this system proved powerful and flexible
- can be extended - broader usage
- programs must be interrupted in other situations as well
 - not only for communication with peripherals
- especially useful for the operating system

Types of Interrupts

- *hardware* - generated by peripherals
- *traps* - generated by the processor itself
 - indicate an abnormal situation
 - example: division by 0
- *software* - generated by programs
 - used to request certain services to the operating system