

V. The Operating System

The Role of the Operating System (1)

- program that manages the computer
- links the hardware to the applications
- provides various services to the applications
- supervises the way applications work
 - can take control when errors occur

The Role of the Operating System (2)

- needs hardware support to carry out its duties
 - the most important - the interrupt system
- main components
 - kernel
 - drivers

V.1. The Kernel

The Kernel

- mostly independent from the hardware structure it works with
- "the brain" of the operating system
- manages the computing resources - hardware and software
 - proper working
 - fair allocation between the applications

Processor's Operating Modes

- user mode
 - restricted
 - access to memory - only certain zones
 - access to peripheral devices - denied
- kernel mode
 - no restrictions

How Programs Run

- operating system - kernel mode
 - can perform any operation
- applications - user mode
 - cannot perform certain actions
 - must ask the kernel to perform those actions for them

Switching between the Two Modes

- through the interrupt system
- user \rightarrow kernel
 - a software interrupt is called
 - an exception (error) occurs
- kernel \rightarrow user
 - return from the interrupt handling routine

Consequences

- no application code can run in kernel mode
- advantage: the errors of an application do not affect other programs
 - other applications
 - the operating system itself
- drawback: loss of performance

The Structure of the Kernel

- not a unitary program
- more like a collection of cooperating routines
- main functions
 - process management
 - memory management
 - file system management
 - inter-process communication

V.2. System Calls

System Calls

- requests made by applications to the kernel
- actions that applications may not execute themselves
 - can only be performed in processor's kernel mode
 - reason - system safety
- achieved through software interrupts
- similar to function calls

The Phases of a System Call (1)

1. the program lays the parameters of the system call into a certain zone in memory
2. a software interrupt is generated
 - processor switches to kernel mode
3. the requested service is identified and the corresponding routine is called

The Phases of a System Call (2)

4. the routine takes the parameters and checks them
 - if there are errors - the call fails
5. if there are no errors - the requested action is performed
6. the routine ends - results are laid into a memory zone that is accessible to the requesting application

The Phases of a System Call (3)

7. processor switches back to user mode
8. program execution is resumed from the point it had been interrupted
 - using the information memorized when interrupt was generated
9. the program can now take the results of the call

System Calls - Conclusions

- communication between the application and the kernel
 - lay the parameters
 - take the results
- critical actions are performed safely
- high time consuming
- make calls seldom - work with buffers

How We Use Buffers

Example - the *printf* function

- formats the text, then sends it to screen
- no direct access to hardware
 - uses a system call
 - *write* (the Linux case)
- in fact, *printf* keeps the formatted text into a local memory zone (buffer)
 - a system call is made when the buffer is full

V.3. Drivers

What Are Drivers?

- program modules which manage communication to peripheral devices
 - specialized - a specific driver for each peripheral device
- part of the operating system
 - direct access to hardware
 - run in processor's kernel mode

Usage

- drivers are not part of the kernel
 - but are controlled by the kernel
- used by the hardware interrupt handling routines
- change a peripheral → change the driver
 - modular structure
 - no need to re-install the entire operating system

V.4. Process Management

Processes (1)

- one can launch more programs at the same time (multitasking)
- parallelism is not real
 - unless the system has more processors
 - otherwise - concurrency
- a program may be split into more sequences of instructions - *processes*
 - can be executed in parallel or concurrently

Processes (2)

- the operating system works with processes
 - not with programs
- when launched, a program consists in a single process
 - it can create other processes
 - which can create other processes, and so on
- at each moment, a processor can execute the instructions of a single process

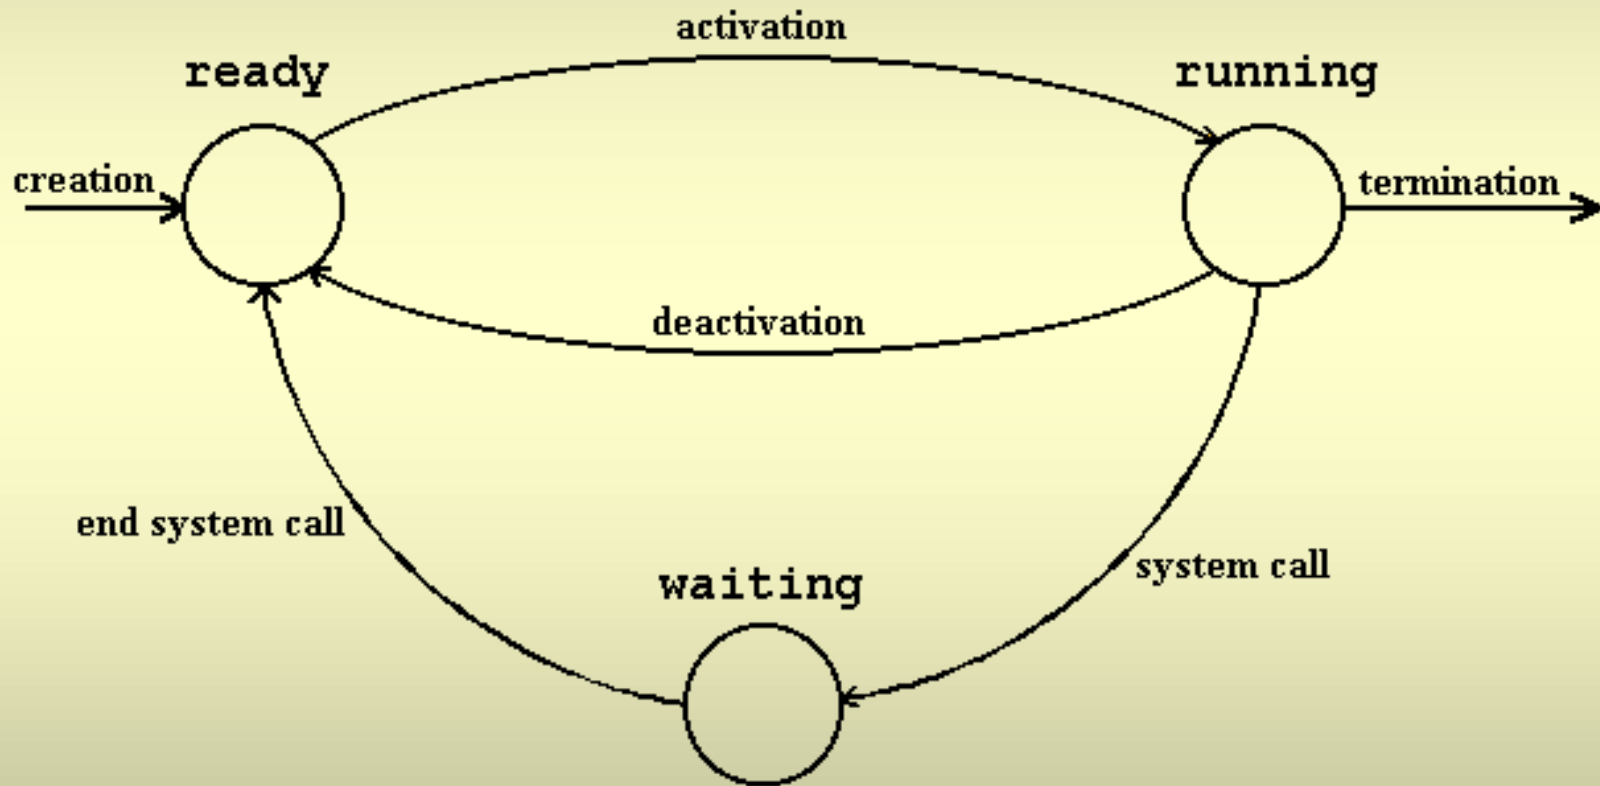
Processes (3)

- each process has its own memory zones (code, data, stack, ...)
 - separated from the other processes' zones
- when a new process is created, proper memory space is allocated to it
- when a process is finished, the memory occupied by that process is released
 - even if the program as a whole continues

States of a Process (1)

- running
 - its instructions are executed by the processor
- ready
 - waits to be executed by the processor
- waiting
 - waits for the completion of a system call
 - currently, it does not compete for being assigned to the processor

States of a Process (2)



States of a Process (3)

- the running process leaves this state
 - when it ends
 - normally or because of an error
 - when it makes a system call (→ waiting)
 - when its instructions have been executed long enough and it is time for another process to be executed (deactivation)

Multitasking

- non-preemptive
 - does not allow the deactivation of a process
 - a process can stop executing only in one of the other situations described before
 - drawback - programming errors may block the other processes (e.g., an infinite loop)
- preemptive

Deactivation

- how does the operating system know how long has a process been executed?
- a way of measuring the time is required
- real-time clock
 - peripheral device
 - generates interrupt requests periodically

Threads (1)

- a process can be split into more execution threads
 - usually, a thread consists in the execution of a function (part of the process' code)
- threads share the resources of the process (memory, open files, etc.)
 - simpler communication - global variables
 - higher risk of (undesired) interference

Threads (2)

- when a process is scheduled to the processor, one of its threads will execute
 - so some kind of scheduling is needed here too
- who performs the scheduling?
- alternatives
 - the operating system (rather seldom)
 - the application - through dedicated library functions

V.5. Memory Management

Memory Management

Functions

- allocates memory zones to the applications
- prevents the interferences between applications from occurring
- detects and stops illegal accesses


Fundamental Problem

- multiple applications running → separate memory zones
- each application → certain memory zones; which are these zones?
 - depend on memory configuration at that moment
 - cannot be known at compiling time

Solution

- two kinds of addresses
 - virtual - which the application believes to access
 - physical - really accessed by the processor
- the correspondence between virtual and physical addresses - managed by the operating system

Virtual and Physical Addresses (1)



Virtual and Physical Addresses (2)

Managing virtual and physical addresses

- 2 different methods
 - segmentation
 - pagination
- can also be used together
- dedicated component of the processor - MMU (Memory Management Unit)

V.5.1. Memory Segmentation

Basic Principle (1)

- segment - contiguous memory zone
- contains information of the same kind (code, data, etc.)
- visible to the programmer
- the address of a location - made of 2 parts
 - the start address of the segment
 - the offset within the segment

Basic Principle (2)

- upon different executions of the program, a segment starts at different addresses
- the effect on location addresses
 - the start address of the segment must be updated
 - the offset - not modified
- the problem is only partially solved
 - we want the address not to be modified at all

Descriptors (1)

- segment descriptor - data structure for managing a segment
- the information it keeps
 - start address
 - size
 - access rights
 - etc.

Descriptors (2)

- descriptors - grouped into a table
- access to a segment - based on the index in the descriptor table (selector)
- virtual address - 2 components
 - the index in the descriptor table
 - the offset within the segment
- physical address = start address of the segment + offset

Descriptors (3)

- upon different executions of the program, a segment starts at different addresses
- the effect on location addresses
 - none
 - must only modify the start address of the segment in the descriptor
 - only once (when the segment is loaded into memory)
 - performed by the operating system

Memory Access (1)

- the program requests the virtual address
- the segment descriptor is identified
- the access rights are checked
 - insufficient rights - a trap is generated
- the offset is checked
 - if the offset exceeds the segment size - a trap is generated

Memory Access (2)

- if an error has occurred on previous steps
 - the service routine of the trap ends the program
- if no error has occurred
 - compute the physical address (segment start address + offset)
 - access the location at the computed address

Illustration (1)

Descriptor table (simplified)

Index	Start address	Size
0	65000	43000
1	211000	15500
2	20000	30000
3	155000	49000
4	250000	35000

Illustration (2)

Example 1:

```
mov byte ptr ds:[eax], 25
```

- $ds = 3 \rightarrow$ segment start address = 155000
- $eax = 27348 < 49000$
 - valid offset (does not exceed the segment size)
- physical address: $155000 + 27348 = 182348$

Illustration (3)

Example 2:

```
add dword ptr ss:[ebp], 4
```

- $ss = 1 \rightarrow$ segment start address = 211000
- $ebp = 19370 > 15500$
 - invalid offset (exceeds the segment size)
 - error \rightarrow generate trap

Intel Case (1)

3 descriptor tables

- global (GDT - Global Descriptor Table)
 - accessible to all processes
- local (LDT - Local Descriptor Table)
 - specific to each process
- interrupts (IDT - Interrupt Descriptor Table)
 - not directly accessible to the applications

Intel Case (2)

Segments - accessed through the selectors

Selector structure (16 bits)

- first 13 bits - the index in the descriptor table
 - at most 8192 descriptors/table
- 1 bit - the table used (global/local)
- last 2 bits - privilege level
 - 0 - highest, 3 - lowest

Intel Case (3)

31	24				19	16	15	14	13	12	8	7	0
BASE 31-24				G	D / B	0	A V L	LIMIT 19-16		P	DPL	TYPE	BASE 23-16
BASE 15-0									LIMIT 15-0				

Intel Case (4)

- privilege levels of 3 entities are considered
 1. CPL (*Current Privilege Level*)
 - level of the process - kept by the processor
 2. RPL (*Requested Privilege Level*)
 - the requested level - taken from the selector
 3. DPL (*Descriptor Privilege Level*)
 - the accessed segment's level - from the descriptor

Intel Case (5)

- relations between these privilege levels decide if access is allowed
- the condition for allowing the access:
 $CPL \leq DPL$ and $RPL \leq DPL$
(simultaneously)
- any other situation indicates an attempted access to a level too high
 - a trap is generated