

Introduction to Distributed Systems

C.Vidraşcu

<https://profs.info.uaic.ro/~vidrascu>

Agenda

- Flynn's taxonomy
- Distributed systems types
- Design issues for DS
- History of DS

Flynn's taxonomy (1966) (1)

- The processor (CPU) has 2 functional parts:
 - the control unit : decides what operations need to be done
 - the processing unit : executes those operations
- A control unit may command one or several processing units


Flynn's taxonomy (2)

- SISD
 - Single Instruction, Single Data
- SIMD
 - Single Instruction, Multiple Data
- MISD
 - Multiple Instruction, Single Data
- MIMD
 - Multiple Instruction, Multiple Data

SISD (1)

- only one control unit
- only one processing unit
- von Neumann architecture
- it works completely sequentially:

the execution of the next instruction only starts after finishing the execution of the previous one



SISD (2)

- Examples:

UNIVAC1



IBM 360



CRAY1



CDC 7600



PDP1




Dell Laptop




SIMD (1)

- one control unit and several processing units
- the same instruction is executed in parallel on different data
- implementations:
 - vector pipelines (e.g. IBM 9000)
 - processor arrays (e.g. ILLIAC IV)



SIMD (2)



SIMD (3)

- Examples:


ILLIAC IV



MasPar



Cell Processor (GPU)




Cray X-MP



Cray Y-MP



Thinking Machines CM-2




SIMD (4)

- not all operations can be done in parallel
- in practice, combinations SISD-SIMD
 - Intel: MMX, SSE
- the problem is caused by the conditional instructions
 - instruction streams diverge because of them
 - solution: disabling of processing units with incorrect data
 - drawback: loss of performance


MISD

- several control units and only one processing unit
- different operations executed on the same data, in parallel
- implementation - ???
- pipeline systems
 - superscalar processor computers (e.g. CDC 6600)
- redundant systems

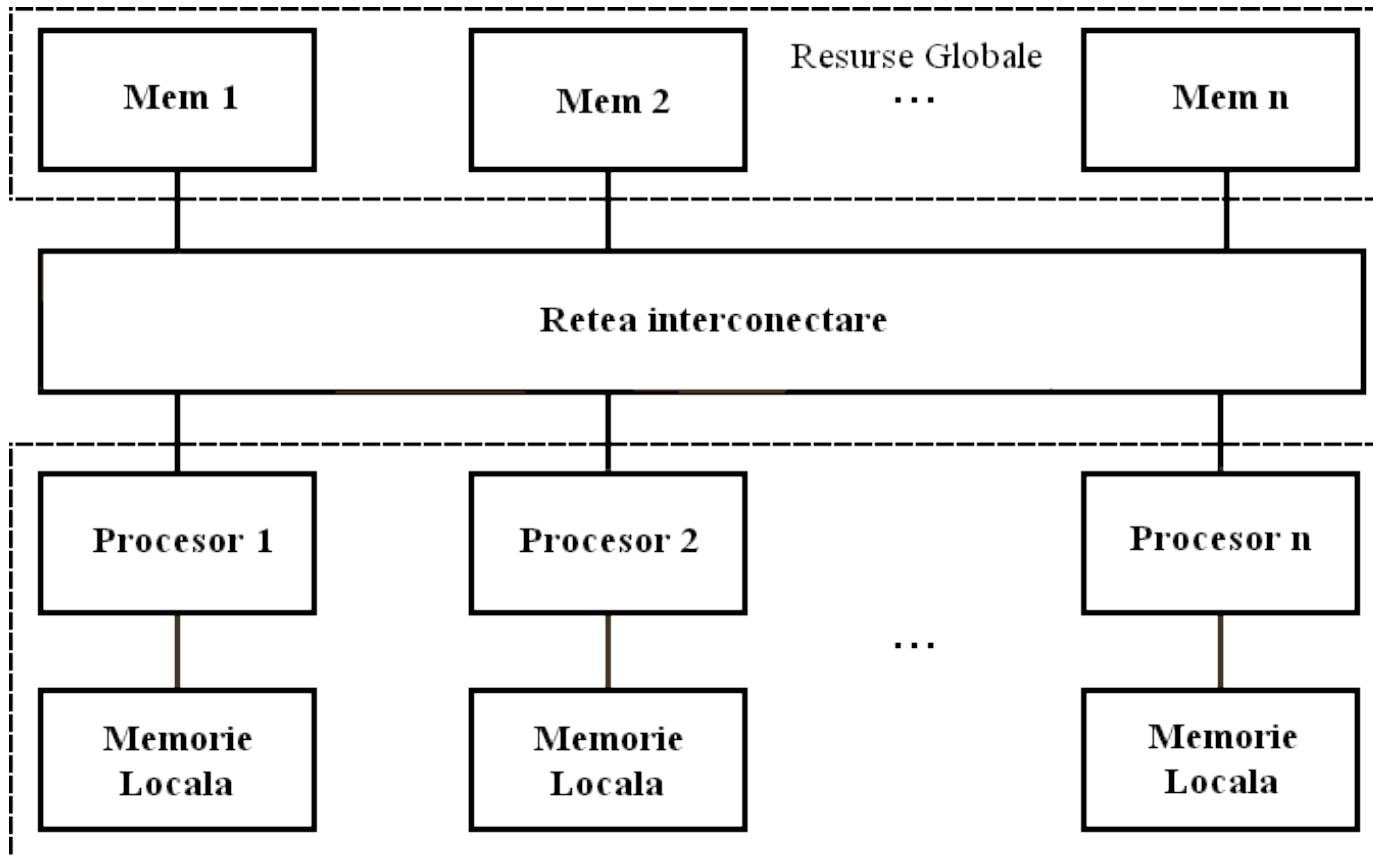


MIMD (1)

- several control units and several processing units
- more than one instruction streams, executed:
 - independently
 - cooperatively
- distributed systems are MIMD



MIMD (2)



MIMD (3)

- Examples:

IBM POWER5



HP/Compaq Alphaserwer



Intel IA32



AMD Opteron



Cray XT3



IBM BG/L



MIMD - classification

- computers can have shared memory or not
 - computers with shared memory - multiprocessors
 - all processors have the same address space for accessing the memory
 - generally speaking, are called parallel systems
 - computers without shared memory - multicomputers
 - each processor has its own address space
 - generally speaking, are called distributed systems
- based on the interconnection topology:
 - bus vs. switched

Interconnection

- *tightly coupled* systems
 - fast transfer of messages between the processing units
 - transfer data rate - high
- *loosely coupled* systems
 - slow transfer, low data rate
- this classification can be done both for software, as well as for hardware

Distributed systems types

- different combinations are possible based on the hardware/software interconnection types
- *parallel systems*: systems which are tightly coupled in hardware
- *distributed systems* are usually loosely coupled in hardware
- variants of *distributed systems*:
 - network operating systems
 - "true" distributed systems

Parallel systems

- tightly coupled software, running on tightly coupled hardware
- hardware: a single multiprocessor computer (i.e. a computer with several physical CPUs or, nowadays, with several cores in CPU)
- software: e.g. any time-sharing operating system for SMPs (UNIX, Windows, etc.)

Network operating systems (1)

- loosely coupled software, running on loosely coupled hardware
- independent computers (workstations), connected in a LAN network
- each workstation has its own operating system
 - the same operating system on all workstations
 - or
 - not, i.e. different OSes, or different versions of it

Network operating systems (2)

- from one workstation tasks can be executed on other workstations:
 - running some commands/programs remotely
 - file transfer
- more complex communications mechanisms
 - ex.: file server
- there has to be an agreement between all workstations about the syntax and the meaning of the messages transferred between them
 - i.e. protocols for communication

"True" distributed systems (1)

- tightly coupled software, running on loosely coupled hardware
- the users see the distributed system as a single classical computer (*single-system image*)
 - what kind of user? end-users, but also programmers
 - end-users point of view: the running applications see the system's resources in a unified manner, no matter which machines they run on
 - programmers p.o.v.: system calls hide the distributed details

"True" distributed systems (2)

Characteristics of a distributed OS:

- Unified mechanism for interprocess communication
- Uniform management of processes
 - not acceptable to have different scheduling policies on different individual machines
- Unified file system
- Conclusion: all machines are running the same OS kernel
- Plus: a higher-level collection of *system management components* that coordinate the machines' activities

Design issues

- Distributed systems have (also) requirements different from those of classical SISD computers
- When we design a distributed system, what issues we have to deal with?
 - transparency
 - reliability
 - performance
 - scalability

Transparency

- at the end-user level
 - when running a command, the system hides from the user on how many (and which) computers that command will run
- at the programmer level
 - the system calls interfaces hide the distributed structure of the system

Aspects of transparency (1)

- location transparency
 - hardware and software resources can be located anywhere in the sistem
 - they can be accessed in the same way from any machine
 - examples of resources: processors, files, peripheral devices, etc.

Aspects of transparency (2)

- migration transparency
 - extends the location transparency for moving resources
 - if a resource migrates from one machine to another, must be possible to access it using the same name as before the migration
 - this issue is important for migrating resources between servers, without downtime

Aspects of transparency (3)

- replication transparency
 - The OS can make copies (duplicates) of resources without the users knowing about it
 - example: a couple of file servers, which works together to provide a distributed file-system
 - each server stores a subset of the files
 - when a client makes a request to a server, if the request implies a file stored somewhere else, the server forwards the request and sends back to the client the file received, but can keep a local copy of it for future accesses
 - goal: for performance

Aspects of transparency (4)

- concurrency transparency
 - refers to the concurrent accesses at some resources made by several users/programs
 - the users will not notice the existence of other users
 - the accesses requested in parallel are solved by the OS sequentially

Aspects of transparency (5)

- parallelism transparency
 - goal: the programmers write applications as for classical uniprocessor computers, but the compiling and/or runtime environment takes advantage of the parallel computing capabilities of the system
 - in parallel computing theory this is called: implicit parallelism
 - It is hard to achieve in the current state of IT

Reliability

- A distributed system can have many computing nodes
 - it is very likely that some components goes down (i.e. malfunction)
 - availability – the periods of time in which the system is non-functional to be as short as possible
 - fault-tolerance – the system must be capable to function even when some components are down

Performance

- the reason for developing parallel and distributed system was to have an increased performance (i.e. computing power)
- these kind of systems are not justified economically for running simple tasks, which can be solved by classical computers
- the limiting factor of performance in distributed systems is communication
 - hardware and software
 - *fine-grained parallelism vs. coarse-grained parallelism*

Scalability

- describes how the performance of a parallel or distributed system is increased when the number of CPUs/machines is increased
- preferable – linear increase
- linear increase is hard to get
 - some applications are inherently sequential
 - communication between nodes – bottleneck
 - centralized components or data tables – bottleneck
 - centralized vs. distributed algorithms

History of distributed systems

- 1954: DYSEAC, a general-purpose synchronous computer, with distributed control
- 1970s-1980s: important research advances in distributed computing, but implementations had modest commercial success
- 1990s: *clusters* of commodity hardware (PCs)
 - improved performance and availability over that of a single PC
- 2000s: *grids* – "super virtual computers" & Virtual Org.
 - a collection of computer resources from multiple locations (organisations/enterprises, etc.) to reach a common goal of them
- After 2007: *clouds* – grids with users from outside the VO
 - service models: SaaS, PaaS, IaaS
 - deployment models: public, private, hybrid

- **Bibliografie obligatorie**

capitolele despre *sisteme distribuite* din

- Silberschatz : “*Operating System Concepts*”

(cap.19 din [OSC10])

sau

- Tanenbaum : “*Modern Operating Systems*”

(cap.8 din [MOS4])

- **Bibliografie recomandată**

- Andrew S. Tanenbaum : “*Introduction to Distributed Systems*”,
Prentice Hall, 1995